

YEFAN ZHOU

yefan.zhou.gr@dartmouth.edu | Homepage [↗](#) | Google Scholar [↗](#) | Linkedin [↗](#) | Hanover, NH, 03755 | 510-809-5378

EDUCATION

Dartmouth College

Ph.D. student in Computer Science

Hanover, NH

Sep. 2023 – present

Advisor: Prof. Yaoqing Yang

Research Area: Efficiency and transparency of ML, LLM pruning/fine-tuning/mixture-of-expert, Model diagnostic

University of California, Berkeley

Berkeley, CA

M.Eng in Electrical Engineering and Computer Science; Major GPA: 4.0/4.0

Aug. 2021 – Dec. 2022

Advisor: Prof. Michael Mahoney

Research Area: Pruning for model efficiency

University of California, Berkeley

Berkeley, CA

Exchange Student; GPA: 4.0/4.0

Jan. 2019 – May. 2019

Southeast University

China

B.Eng in Information Engineering; GPA: 3.7/4

Aug. 2016 – Jun. 2020

PUBLICATION

Selected first-author paper:

- {H. Lu*, **Y. Zhou***}, S. Liu, Z. Wang, M. W. Mahoney, Y. Yang “AlphaPruning: Using Heavy-Tailed Self Regularization Theory for Improved Layer-wise Pruning of Large Language Models” (NeurIPS 2024)
- {H. Lu*, X. Liu*, **Y. Zhou***, Q. Li*}, H. Yang, Y. Yan, K. Keutzer, M. W. Mahoney, Y. Yang “Sharpness-diversity tradeoff: improving flat ensembles with SharpBalance” (NeurIPS 2024)
- {**Y. Zhou***, J. Chen*}, Q. Cao, K. Schürholt, Y. Yang “MD tree: a model-diagnostic tree grown on loss landscape” (ICML 2024)
- {**Y. Zhou***, T. Pang*}, K. Liu, C. H. Martin, M. W. Mahoney, Y. Yang “Temperature Balancing, Layer-wise Weight Analysis, and Neural Network Training” (NeurIPS 2023 Spotlight)
- **Y. Zhou**, Y. Yang, A. Chang, M. W. Mahoney “A Three-regime model of Network Pruning” (ICML 2023)
- **Y. Zhou**, Y. Shen, Y. Yan, C. Feng, Y. Yang “A Dataset-Dispersion Perspective on Reconstruction Versus Recognition in Single-View 3D Reconstruction Networks” *2021 International Conference on 3D Vision (3DV 2021)*

Collaborating or advising paper:

- {Z. Liu*, Y. Hu*}, T. Pang, **Y. Zhou**, P. Ren, Y. Yang “Model Balancing Helps Low-data Training and Fine-tuning” (EMNLP 2024 main)
- P. Qing, C. Gao, **Y. Zhou**, X. Diao, Y. Yang, S. Vosoughi “AlphaExpert: Assigning LoRA Experts Based on Layer Training Quality” (EMNLP 2024 main)
- X. Zhu, **Y. Zhou**, Y. Fan, J. Chen, M. Tomizuka “Learn to Grasp with Less Supervision: A Data-Efficient Maximum Likelihood Grasp Sampling Loss” *2022 International Conference on Robotics and Automation (ICRA 2022)*
- K. Schürholt, L. Meynert, **Y. Zhou**, Y. Yang, D. Borth “A Model Zoo on Phase Transitions in Neural Networks” (Preprint)

PROFESSIONAL EXPERIENCE

Research Engineer, International Computer Science Institute

Berkeley, CA

supervised by Prof. Michael Mahoney

Jan. 2023 – Jun. 2023

- Researched efficient optimization method for deep neural network.
- Researched ensembling methods for improving the OOD robustness of CV models.
- Developed backdoor detection methods to enhance AI model safety.

Graduate Research Assistant, Sky Computing Lab (RISELab), UC Berkeley

Berkeley, CA

advised by Prof. Michael Mahoney

Aug. 2021 – Dec. 2022

- Researched neural network pruning for CNNs and Transformers.

SERVICES AND AWARD

Reviewers: ICLR 2025-2024, NeurIPS 2023, AAAI 2024, ICML 2024, CVPR 2024, CPAL 2024, IROS 2022, TMLR

Talk

- * Invited talk at AI-TIME, “Phase transition, loss landscape and model diagnostics”, Jan., 2024.
- * Invited talk at UC Berkeley/ICSI TrojAI onsite, “Layer-wise Weight Analysis, and Neural Network Training” Oct. 2023
- * Invited talk at UC Berkeley/ICSI TrojAI onsite, “A Three-regime model of Network Pruning” Mar. 2023

Award: ICML 2024 Scholar Award, NeurIPS 2023 Scholar Award

Teaching (Head TAs): CS70: Foundations of Applied Computer Science (Dartmouth College Spring 2024)

SKILLS

Programming Language: Python, Java, C/C++, CUDA, SQL, MATLAB

Developer Tools: PyTorch, Ubuntu, MuJoCo, ROS, PyBullet, Slurm, PyRender, Open3D