Yefan Zhou

yefan.zhou.gr@dartmouth.edu | Homepage | Google Scholar | Linkedin | Hanover, NH, 03755 | 510-809-5378

Education

Dartmouth College — Hanover, NH
Ph.D. candidate in Computer Science — *Sep. 2023 – present*
Advisor: Prof. Yaoqing Yang and Prof. Michael Mahoney
Research Area: Efficiency and transparency of ML: LLM pruning/fine-tuning, Model diagnostic, Data selection

University of California, Berkeley — Berkeley, CA
Master in Electrical Engineering and Computer Science; Major GPA: 4.0/4.0 — *Aug. 2021 – Dec. 2022*
Advisor: Prof. Michael Mahoney
Research Area: Pruning for model efficiency

University of California, Berkeley — Berkeley, CA
Exchange Student; GPA: 4.0/4.0 — *Jan. 2019 – May. 2019*

Southeast University — China
B.Eng in Information Engineering; GPA: 3.7/4 — *Aug. 2016 – Jun. 2020*

Research Interests

- **LLM pruning**[1, 5]: Proposed an approach to automatically allocate compression factors (such as sparsity or quantization precision) to transformer layers, applicable to LLaMA, OPT, Mistral, ViT, and ConvNext. Pruning LLaMA-7B to 80% sparsity leads to 3.06× end-to-end speedup on CPUs tested on the DeepSparse kernel. It also improves the structured/semi-structured pruning and mix-precision quantization.
- **Model training optimization**[4, 7, 8]: Proposed an adaptive learning rate optimizer to balance the layer qualities of models for model training and LLM fine-tuning. It consistently improves tasks of image classification, question answering, and Neural PDE solving. It improves the RoBERTa-base fine-tuned on the SST2 dataset and increases the accuracy of LLaMA-7B on the ScienceQA dataset by 1.97%.
- **Training data selection**[2]: Proposed a method for dynamically selecting data for Ensembling training to enhance the robustness of CV models in out-of-distribution image classification.
- **Transparency and Explainability**[3]: Proposed a post-training diagnostic method to identify pre-trained model failures (e.g., insufficient model size) without access to the training configuration, streamlining hyperparameter tuning.

[·] refers to publications listed below.

Publication

Selected first-author paper:

1. {H. Lu*, **Y. Zhou***}, S. Liu, Z. Wang, M. W Mahoney, Y. Yang "AlphaPruning: Using Heavy-Tailed Self Regularization Theory for Improved Layer-wise Pruning of Large Language Models" (**NeurIPS 2024**)
   LLM pruning · Efficient inference
2. {H. Lu*, X. Liu*, **Y. Zhou***, Q. Li*}, H. Yang, Y. Yan, K. Keutzer, M. W. Mahoney, Y. Yang "Sharpness-diversity tradeoff: improving flat ensembles with SharpBalance" (**NeurIPS 2024**)
   Training data selection · Ensembling · Out-of-distribution
3. {**Y. Zhou***, J. Chen*}, Q. Cao, K. Schürholt, Y. Yang "MD tree: a model-diagnostic tree grown on loss landscape" (**ICML 2024**)
   Model selection · Scaling law · Hyperparameter tuning
4. {**Y. Zhou***, T. Pang*}, K. Liu, C. H Martin, M. W Mahoney, Y. Yang "Temperature Balancing, Layer-wise Weight Analysis, and Neural Network Training" (**NeurIPS 2023 Spotlight**)
   NN optimizer · Efficient training · Layer quality analysis
5. **Y. Zhou**, Y. Yang, A. Chang, M. W Mahoney "A Three-regime model of Network Pruning" (**ICML 2023**)
   NN pruning · Model selection · Losslandscape analysis
6. **Y. Zhou**, Y. Shen, Y. Yan, C. Feng, Y. Yang "A Dataset-Dispersion Perspective on Reconstruction Versus Recognition in Single-View 3D Reconstruction Networks" *2021 International Conference on 3D Vision* (**3DV 2021**) 3D reconstruction · 2D-to-3D

Collaborating or advising paper:

7. {Z. Liu*, Y. Hu*}, T. Pang, **Y. Zhou**, P. Ren, Y. Yang "Model Balancing Helps Low-data Training and Fine-tuning" (**EMNLP 2024 main Oral**)
   LLM fine-tuning Low-resource training
8. P. Qing, C. Gao, **Y. Zhou**, X. Diao, Y. Yang, S. Vosoughi "AlphaExpert: Assigning LoRA Experts Based on Layer Training Quality" (**EMNLP 2024 main**)
   LLM fine-tuning Mixture-of-expert
9. X. Zhu, **Y. Zhou**, Y. Fan, J. Chen, M. Tomizuka "Learn to Grasp with Less Supervision: A Data-Efficient Maximum Likelihood Grasp Sampling Loss" *2022 International Conference on Robotics and Automation* (**ICRA 2022**) Few-shot learning
10. K. Schürholt, L. Meynent, **Y. Zhou**, Y. Yang, D. Borth "A Model Zoo on Phase Transitions in Neural Networks" (Preprint)

Professional Experience

Research Engineer, International Computer Science Institute Berkeley, CA
supervised by **Prof. Michael Mahoney** *Jan. 2023 – Jun. 2023*

- Researched efficient optimization method for deep neural network.
- Researched ensembling methods for improving the OOD robustness of CV models.
- Developed backdoor detection methods to enhance AI model safety.

Graduate Research Assistant, Sky Computing Lab (RISELab), UC Berkeley Berkeley, CA
advised by **Prof. Michael Mahoney** *Aug. 2021 – Dec. 2022*

- Researched neural network pruning for CNNs and Transformers.

Services and Award

Reviewers: ICLR 2025-2024, NeurIPS 2023, AAAI 2024, ICML 2024, CVPR 2024, CPAL 2024, IROS 2022, TMLR

Talk

* Invited talk at AI-TIME, "Phase transition, loss landscape and model diagnostics", Jan., 2024.
* Invited talk at UC Berkeley/ICSI TrojAI onsite, "Layer-wise Weight Analysis, and Neural Network Training" Oct. 2023
* Invited talk at UC Berkeley/ICSI TrojAI onsite, "A Three-regime model of Network Pruning" Mar. 2023

Award: ICML 2024 Scholar Award, NeurIPS 2023 Scholar Award

Teaching (Head TAs): CS70: Foundations of Applied Computer Science (Dartmouth College Spring 2024)

Skills

Programming Language: Python, Java, SQL, MATLAB

Developer Tools: PyTorch, Ubuntu, MujoCo, ROS, PyBullet, Slurm, PyRender, Open3D